

Problem Statement and Goals

Software Engineering

Team #1, Sanskrit Ciphers
Omar El Aref
Dylan Garner
Muhammad Umar Khan
Aswin Kuganesan
Yousef Shahin

Table 1: Revision History

Date	Developer(s)	Change
09/13/2025	Omar El Aref(s)	Added the problem statement, Inputs and Outputs
09/14/2025	Omar El Aref(s)	Added Stakeholders, Environment, Goals, Stretch Goals and Extras
09/22/2025	Omar El Aref(s)	Added Reflection

1 Problem Statement

1.1 Problem

The textual history of Indian Buddhism is fragmented across thousands of manuscript folios, preserved only in partial, damaged, or scattered forms. Traditionally, scholars reconstruct these texts manually through palaeographic study, transcription, and content comparison; a slow and error-prone process. While some existing computational tools exist for pattern recognition and Optical Pattern Recognition (OCR) none can tackle this problem as they are not designed for irregular, damaged, or arbitrarily oriented manuscript fragments.

The lack of computational tools available for this problem significantly limits progress in reconstructing Buddhist textual history. Scholars need a tool that automates the detection, matching, and transcription of manuscript fragments, thereby reducing manual effort and time; enabling large-scale reconstruction.

1.2 Inputs and Outputs

- **Inputs:**

- High quality images of manuscript fragments (approximately 21,000 images from collections)
- Metadata (collection identifiers, orientation, partial transcriptions if available)

- **Outputs:**

- Probabilistic matches between fragments based on shape/edge/damage features (i.e., list of fragments most likely related to an input fragment image)
- Enhanced Metadata:
 - * Suggested transcriptions of fragment text
- Searchable/sortable database containing fragment attributes and relationships

1.3 Stakeholders

- **Primary Stakeholders:**

- Scholars of Buddhist textual history (religious studies, philology, palaeography)
- Supervisors and domain experts (e.g., Dr. Shayne Clarke, McMaster Religious Studies)

- **Secondary Stakeholders:**

- Computer science researchers in AI, image processing
- Humanities researchers studying textual transmission and manuscript culture

- **End Users:**

- Academic researchers using the software to assemble fragments
- Graduate students seeking computational support in philological research

1.4 Environment

- **Hardware:**

- Development laptops/workstations with GPU support for ML tasks

- **Software:**

- VScode
- Coding Libraries
- Database
- GitHub repository for version control and CI/CD

2 Goals

- Develop a tool that:
 - Detects edges, shapes, and damage patterns in fragments.
 - Matches fragments based on similarity measures (probabilistic scoring).
 - Identifies palaeographic script of fragments.
 - Performs preliminary transcription using OCR tuned to Sanskrit scripts.
 - Builds a searchable database linking fragments with metadata and probable matches.
 - Provides a user interface for scholars to view suggested fragment matches and confirm/annotate them.

3 Stretch Goals

- Expand support to Tibetan and Chinese manuscript fragments.
- Incorporate semantic content matching with other fragments.
- Improve transcription accuracy with AI-assisted error correction.

4 Extras

- **User Documentation:** Write user-friendly guides for scholars with limited technical expertise.
- **Use Case Video:** Make a Video on how to use the tool so that we can minimize the learning curve for the intended user

Appendix — Reflection

The purpose of reflection questions is to give you a chance to assess your own learning and that of your group as a whole, and to find ways to improve in the future. Reflection is an important part of the learning process. Reflection is also an essential component of a successful software development process.

Reflections are most interesting and useful when they're honest, even if the stories they tell are imperfect. You will be marked based on your depth of thought and analysis, and not based on the content of the reflections themselves. Thus, for full marks we encourage you to answer openly and honestly and to avoid simply writing "what you think the evaluator wants to hear."

Please answer the following questions. Some questions can be answered on the team level, but where appropriate, each team member should write their own response:

1. What went well while writing this deliverable?
2. What pain points did you experience during this deliverable, and how did you resolve them?
3. How did you and your team adjust the scope of your goals to ensure they are suitable for a Capstone project (not overly ambitious but also of appropriate complexity for a senior design project)?

What went well the most would be that everyone was on the same page for this project and so coming up with what we wanted to accomplish was pretty easy and we didn't really have many disagreements. This made it easy to focus on the task of writing the problem statement and goals rather than having to worry about different opinions on the team. This project was also one of the projects that was on the list of projects pdf given to us so it was easy to figure out the problem statement and goals since we didn't really have much to come up with anything on our own.

One challenge was avoiding either too much technical detail or too much abstraction. We really had to focus on what we wanted to convey and add too much detail as to not constrain our selves but also not have too little detail that it isn't clear as to what we are doing. We initially struggled to find the right balance between scholarly needs (manuscript context) and technical specifications (ML algorithms, environment). We resolved this by starting broad, then refining with feedback and checking against the POC and problem statement checklists. That way we made sure that our POC and problem statement were in line with each other. Another pain point was uncertainty about which machine learning techniques would realistically be feasible; to address this, we distinguished between core goals and stretch goals to avoid overcommitting. This way we also didn't constrain ourselves later down the line when we start implementing the solution.

Initially, we considered a full end-to-end system covering Sanskrit, Tibetan, and Chinese manuscripts. We recognized this was quite ambitious

especially given the time and resources that we had. Instead, we narrowed our core scope to Sanskrit fragments only, focusing on orientation correction, edge/damage-based matching, and preliminary script identification. Transcription and cross-lingual extensions were moved into stretch goals as they are not the core goals that we are trying to achieve with this project. If we are ahead of schedule then they would be great additions to add to the project but again as mentioned they are not the core focus of this project. This adjustment ensures the project is challenging enough, but achievable within the Capstone timeline.